

### Assignment 3

MaxTemperatureMapper.java

```
package MaxMinTemp;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class MaxTemperatureMapper extends
    Mapper<LongWritable, Text, Text, IntWritable> {
    private static final int MISSING = 9999;
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+') { // parseInt doesn't like leading plus
            // signs
            airTemperature = Integer.parseInt(line.substring(88, 92));
        } else {
            airTemperature = Integer.parseInt(line.substring(87, 92));
        }
        String quality = line.substring(92, 93);
        if (airTemperature != MISSING && quality.matches("[01459]")) {
            context.write(new Text(year), new IntWritable(airTemperature));
        }
    }
}
```

MaxTemperatureReducer.java

```
package MaxMinTemp;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class MaxTemperatureReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context)
        throws IOException, InterruptedException {
        int maxValue = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            maxValue = Math.max(maxValue, value.get());
        }
        context.write(key, new IntWritable(maxValue));
    }
}
```

```
MaxTemperatureDriver.java
package MaxMinTemp;
```

```
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
/*This class is responsible for running map reduce job*/
public class MaxTemperatureDriver extends Configured implements Tool{
    public int run(String[] args) throws Exception
    {
        if(args.length !=2) {
            System.err.println("Usage: MaxTemperatureDriver <input path> <outputpath>");
            System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(MaxTemperatureDriver.class);
        job.setJobName("Max Temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(MaxTemperatureMapper.class);
        job.setReducerClass(MaxTemperatureReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true) ? 0:1);
        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }
    public static void main(String[] args) throws Exception {
        MaxTemperatureDriver driver = new MaxTemperatureDriver();
        int exitCode = ToolRunner.run(driver, args);
        System.exit(exitCode);
    }
}
```



```

Map output materialized bytes=1561158
Input split bytes=2109
Combine input records=0
Combine output records=0
Reduce input groups=20
Reduce shuffle bytes=1567158
Reduce input records=142458
Reduce output records=29
Spilled Records=284916
Shuffled Maps =20
Failed Shuffles=0
Merged Map outputs=20
GC time elapsed (ms)=5033
CPU time spent (ms)=19900
Physical memory (bytes) snapshot=6212431872
Virtual memory (bytes) snapshot=41748148224
Total committed heap usage (bytes)=4175953920

```

#### Shuffle Errors

```

BAD ID=0
CONNECTION=0
IO ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

#### File Input Format Counters

```
Bytes Read=19630411
```

#### File Output Format Counters

```
Bytes Written=180
```

```
hduuser@student-OptiPlex-3020:~/TemperatureB4$ SHAD00P_HOME/bin/hdfs dfs -cat /output-temperature/*
```

```
20/01/02 15:10:55 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
1901 317
```

```
1902 244
```

```
1903 289
```

```
1904 256
```

```
1905 283
```

```
1906 294
```

```
1907 283
```

```
1908 289
```

```
1909 278
```

```
1910 294
```

```
1911 306
```

```
1912 322
```

```
1913 300
```

```
1914 333
```

```
1915 234
```

```
1916 276
```

```
1917 317
```

```
1918 322
```

```
1919 378
```

```
1920 294
```

```
hduuser@student-OptiPlex-3020:~/TemperatureB4$
```

ure

Download

File information - part-r-00000

Download

Head the file (first 32K)

Tail the file (last 32K)

Block information

Block 0

Block ID: 1073741881

Block Pool ID: BP-1471511537-127.0.1.1-1571285925503

Generation Stamp: 1057

Size: 180

Availability:

• student-OptiPlex-3020

File contents

1901	317
1902	244
1903	289
1904	256
1905	283
1906	294
1907	283
1908	289

Close